

# Hands-on Introduction to Deep Learning with PyTorch


Training a simple model

Rafael Sarmiento and Rocco Meli

ETHZürich / CSCS

Lugano, February 28<sup>th</sup> - March 1<sup>st</sup> 2024

# How to find a $y_i$ given a $x_i$ that's not in the data?



- We select a linear **model** which is appropriate for our dataset and problem

$$y = mx + n$$

$$y = mx + n$$

- We select a linear **model** which is appropriate for our dataset and problem
- How can we use the data to adjust our **model's parameters** ( $m$  and  $n$ ) to minimize prediction error and ensure accuracy within an acceptable margin?

$$y = mx + n$$

$$L = \frac{1}{N} \sum_i^N (\hat{y}_i - y_i)^2$$

- We select a linear **model** which is appropriate for our dataset and problem
- How can we use the data to adjust our **model's parameters** ( $m$  and  $n$ ) to minimize prediction error and ensure accuracy within an acceptable margin?
- We select a suitable **loss function**
- A loss function quantifies the difference between predicted and actual values, serving as a measure of how well the model performs on a task


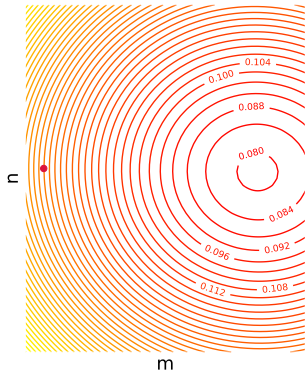
$$y = mx + n$$

$$L = \frac{1}{N} \sum_i^N (\hat{y}_i - y_i)^2$$

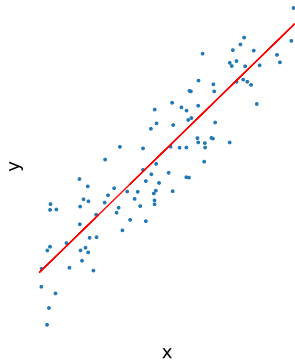
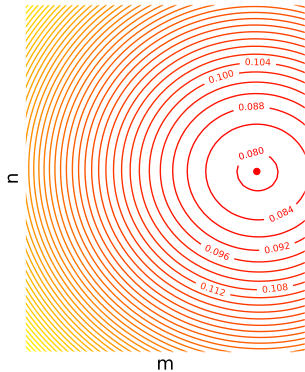
$$L = \frac{1}{N} \sum_i^N (mx_i + n - y_i)^2$$

- We select a linear **model** which is appropriate for our dataset and problem
- How can we use the data to adjust our **model's parameters** ( $m$  and  $n$ ) to minimize prediction error and ensure accuracy within an acceptable margin?
- We select a suitable **loss function**
- A loss function quantifies the difference between predicted and actual values, serving as a measure of how well the model performs on a task
- The loss is a function of both the data and the model parameters. Our objective is to minimize the loss function by adjusting the model parameters (**training the model**)

# We need to choose an **optimizer**




# We need to choose an **optimizer**






# Gradient Descent




- Evaluate the loss function  $L = \frac{1}{N} \sum_i^N l(\hat{y}_i, y_i)$  in the data  $\{x, y\}$  (**forward pass**)

# Gradient Descent



- Evaluate the loss function  $L = \frac{1}{N} \sum_i^N l(\hat{y}_i, y_i)$  in the data  $\{x, y\}$  (**forward pass**)
- Compute the gradients of the loss function with respect to the parameters of the model  $\left. \frac{\partial L}{\partial W} \right|_{\{x, y\}}$  (**backpropagation**)

# Gradient Descent




- Evaluate the loss function  $L = \frac{1}{N} \sum_i^N l(\hat{y}_i, y_i)$  in the data  $\{x, y\}$  (**forward pass**)
- Compute the gradients of the loss function with respect to the parameters of the model  $\left. \frac{\partial L}{\partial W} \right|_{\{x, y\}}$  (**backpropagation**)
- Set a **learning rate** ( $\eta$ ) and update the model parameters  $W_t = W_{t-1} - \eta \left. \frac{\partial L}{\partial W} \right|_{\{x, y\}_{t-1}}$
- Iterate until the minimization of the loss function converges

# Notes on the forward pass and backpropagation

- The computation of the gradients (backpropagation) is done via **Automatic Differentiation**. The chain rule is used to provide exact derivatives of complex functions composed of elementary operations
- During the forward pass, the values of all intermediate operations and its derivatives with respect to the model parameters are stored
- During backpropagation, the gradients of the loss function with respect to the model parameters are computed recursively from the stored values via the chain rule
- PyTorch, TensorFlow, Jax and other deep learning packages provide automatic differentiation


# <Stochastic> Gradient Descent



Gradient  
Descent

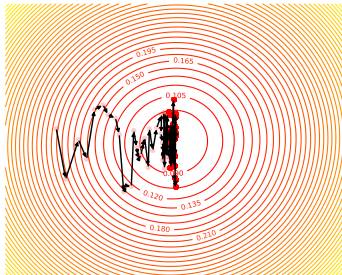
`batch_size = training_set_size`

# <Stochastic> Gradient Descent



Gradient  
Descent


`batch_size = training_set_size`



Stochastic Gradient  
Descent


`batch_size = 1`

# <Stochastic> Gradient Descent




Gradient  
Descent

`batch_size = training_set_size`



Stochastic Gradient  
Descent

`batch_size = 1`



Minibatch Stochastic Gradient  
Descent

$1 < \text{batch\_size} < \text{training\_set\_size}$

# [lab] Simple Stochastic Gradient Descent

Let's run the notebook `sgd/linear_model_sgd.ipynb`. There we use an unidimensional linear model to understand the trajectories of the SGD.

Let's try different batch sizes and learning rates to see how the SGD trajectory changes.



# Summary and Remarks

- Loss function surfaces exhibit high complexity with numerous local minima. Model training primarily focuses on locating suitable local minima rather than striving for the global minimum
- Here we saw an example of **regression** which together with **classification** are the main types of problems in **supervised learning**
- **Unsupervised learning** and **reinforcement learning** are other important categories in machine learning problems
- For simplicity we didn't do any **testing** or **validation** of our model, but that's an important part of training models

Thank you for your attention!